

Business Understanding

In the telecommunications sector, retaining existing customers is more cost-effective than acquiring new ones. High churn rates can lead to significant revenue losses and increased marketing expenses. Understanding the factors that contribute to customer churn enables companies like SyriaTel to implement targeted retention strategies, thereby enhancing customer satisfaction and profitability.

Problem Statement

SyriaTel is experiencing customer attrition, impacting its revenue and market share. The challenge is to develop a predictive model that accurately identifies customers who are likely to churn. By analyzing patterns in customer behavior and demographics, the goal is to proactively address the factors leading to churn and implement effective retention strategies.

Objectives

Regression

Classification

- 1 Predict Churn: Develop a machine learning model to classify customers into churn and non-churn groups.
- 2 Identify Key Drivers: Determine the most influential factors leading to churn.
- 3 Provide Insights: Offer data-driven recommendations for customer retention strategies.
- 4 Improve Customer Retention Efforts: Enable SyriaTel to take proactive actions in reducing churn rates.

Data understanding

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings

warnings.filterwarnings('ignore')
```

```
In [2]: df= pd.read_csv('bigml_59c28831336c6604c800002a.csv')
df.head()
```

Out[2]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34

5 rows × 11 columns



```
In [ ]: df.describe()
```

Out[]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total day charge
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333
mean	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200
std	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50
min	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0
25%	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166
50%	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201
75%	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235
max	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363



```
In [ ]: df.info()
```

```


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   state                                3333 non-null   object
 1   account length                       3333 non-null   int64
 2   area code                           3333 non-null   int64
 3   phone number                         3333 non-null   object
 4   international plan                   3333 non-null   object
 5   voice mail plan                     3333 non-null   object
 6   number vmail messages                3333 non-null   int64
 7   total day minutes                    3333 non-null   float64
 8   total day calls                      3333 non-null   int64
 9   total day charge                     3333 non-null   float64
10  total eve minutes                    3333 non-null   float64
11  total eve calls                      3333 non-null   int64
12  total eve charge                     3333 non-null   float64
13  total night minutes                  3333 non-null   float64
14  total night calls                    3333 non-null   int64
15  total night charge                   3333 non-null   float64
16  total intl minutes                   3333 non-null   float64
17  total intl calls                     3333 non-null   int64
18  total intl charge                    3333 non-null   float64
19  customer service calls               3333 non-null   int64
20  churn                               3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB

```

Data cleaning

```
In [ ]: sns.pairplot(df)
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x7a18c5fa3a90>
```



Correct formats

```
In [ ]: df.head()
```

Out[]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34

5 rows × 21 columns



checking if the data has NAs

In []:

```
# Check NAs
df.isna().mean()*100
```

Out[]:

	0
state	0.0
account length	0.0
area code	0.0
phone number	0.0
international plan	0.0
voice mail plan	0.0
number vmail messages	0.0
total day minutes	0.0
total day calls	0.0
total day charge	0.0
total eve minutes	0.0
total eve calls	0.0
total eve charge	0.0
total night minutes	0.0
total night calls	0.0
total night charge	0.0
total intl minutes	0.0
total intl calls	0.0
total intl charge	0.0
customer service calls	0.0
churn	0.0

dtype: float64

- < 20% - drop rows/fill
- 20 - 50 - fill
- 50% drop column

The data has no missing values

checking if the data has duplicates

Add blockquote

```
In [ ]: df[df.duplicated()].count()
```

```
Out[ ]: 0
```

	0
state	0
account length	0
area code	0
phone number	0
international plan	0
voice mail plan	0
number vmail messages	0
total day minutes	0
total day calls	0
total day charge	0
total eve minutes	0
total eve calls	0
total eve charge	0
total night minutes	0
total night calls	0
total night charge	0
total intl minutes	0
total intl calls	0
total intl charge	0
customer service calls	0
churn	0


dtype: int64

The data has no duplicates

Other cleaning steps

```
In [7]: # Outliers  
  
# Feature engineering
```

```
plt.figure(figsize=(14, 12))
sns.boxplot(data=df)
plt.show()
```



Feature engineering

In []:

EDA

Univariate Analysis


```
In [ ]: # Set up figure for multiple visualizations
plt.figure(figsize=(9, 8))
# Distribution of Account Length for Churn vs. Non-Churned Customers
plt.subplot(2, 2, 1)
sns.histplot(df, x="account length", hue="churn", kde=True, palette="coolwarm", bin
plt.title("Distribution of Account Length by Churn Status")
# Customer Service Calls vs. Churn
plt.subplot(2, 2, 2)
```



```

sns.boxplot(x="churn", y="customer service calls", data=df, palette="coolwarm")
plt.title("Customer Service Calls by Churn Status")
# Total Day Minutes vs. Churn
plt.subplot(2, 2, 3)
sns.boxplot(x="churn", y="total day minutes", data=df, palette="coolwarm")
plt.title("Total Day Minutes by Churn Status")
# International Plan vs. Churn
plt.subplot(2, 2, 4)
sns.countplot(x="international plan", hue="churn", data=df, palette="coolwarm")
plt.title("Churn Rate by International Plan")
plt.tight_layout()
plt.show()

```



```

In [8]: # List of numerical features
numerical_features = ["account length", "total day minutes", "total eve minutes",
                      "total night minutes", "total intl minutes", "customer service


# Set up figure
plt.figure(figsize=(8, 6))

# Histograms & KDE plots for each numerical feature
for i, col in enumerate(numerical_features):
    plt.subplot(3, 2, i + 1) # Create subplots

```

```
sns.histplot(df[col], kde=True, bins=30, color="steelblue")
plt.title(f"Distribution of {col}")

plt.tight_layout()
plt.show()
```



Bivariate analysis

```
In [ ]: # Set up figure for bivariate analysis visualizations
plt.figure(figsize=(9, 6))

# Churn vs. Total Day Minutes
plt.subplot(2, 2, 1)
sns.scatterplot(x=df["total day minutes"], y=df["total day charge"], hue=df["churn"])
plt.title("Total Day Minutes vs. Total Day Charge (Colored by Churn)")

# Churn vs. Customer Service Calls
plt.subplot(2, 2, 2)
sns.boxplot(x="churn", y="customer service calls", data=df, palette="coolwarm")
plt.title("Customer Service Calls by Churn Status")


# Churn vs. Total Intl Minutes & Total Intl Calls
plt.subplot(2, 2, 3)
sns.scatterplot(x=df["total intl minutes"], y=df["total intl calls"], hue=df["churn"])
plt.title("Total Intl Minutes vs. Total Intl Calls (Colored by Churn)")

# Churn vs. Total Evening Usage
```


```
plt.subplot(2, 2, 4)
sns.boxplot(x="churn", y="total eve minutes", data=df, palette="coolwarm")
plt.title("Total Evening Minutes by Churn Status")

plt.tight_layout()
plt.show()
```


Total Day Minutes vs. Total Day Charge (Colored by Churn)




Customer Service Calls by Churn Status



Total Intl Minutes vs. Total Intl Calls (Colored by Churn)



Total Evening Minutes by Churn Status



Observations from Bivariate Analysis

1. Total Day Minutes vs. Total Day Charge:

There's a positive correlation between total day minutes and total day charge, as expected. Customers who churn seem to be spread across various usage levels, without a clear separation in this plot.

2. Customer Service Calls by Churn Status:

Customers who churned tend to have a higher median number of customer service calls. A significant number of churned customers have made 4 or more calls, suggesting that customer service interactions might be a factor in churn.

3. Total Intl Minutes vs. Total Intl Calls:

There's no clear pattern or significant separation between churned and non-churned customers based on international calls and minutes.

4. Total Evening Minutes by Churn Status:

No clear difference in total evening minutes between customers who churned and those who didn't. The distributions overlap substantially.

Overall:


Customer service calls appear to be a potential predictor of churn. Further investigation into the reasons behind these calls may be beneficial. Total day minutes and charges don't show a clear distinction between churned and non-churned customers in this bivariate analysis. International calls and minutes, and evening minutes don't appear strongly related to churn in these visualizations. More in-depth analysis, perhaps with other variables or different visualization techniques, might reveal additional insights.

Multivariate analysis

```
In [9]: # Convert churn to integer type again (to avoid issues with pairplot)
df["churn"] = df["churn"].astype(int)

selected_features = ["total day minutes", "total day charge", "customer service cal

sns.pairplot(df[selected_features], hue="churn", palette="coolwarm", diag_kind="kde
plt.suptitle("Multivariate Analysis: Feature Interactions with Churn", fontsize=14)
plt.show()
```



Key Observations from the Multivariate Analysis

1. **Feature Interactions:** The pairplot helps visualize how different features interact with each other and how these interactions might relate to customer churn. For example, you might observe that customers who churn tend to cluster in specific regions of the "total day minutes" vs. "total day charge" plot. Look for patterns and clusters that are more prevalent among churned customers than non-churned ones.
2. **Combined Effects:** The multivariate analysis allows you to see the combined effects of several features. For example, a customer with high "total day minutes" *and* a high number of "customer service calls" might have a higher probability of churning, something that may not be as apparent from individual bivariate plots.
3. **Marginal Distributions:** The diagonal plots show the distribution of each feature, broken down by churn status. This allows you to compare the distributions of features between churned and non-churned customers. For instance, if the distributions of "customer

service calls" are significantly different (e.g., churned customers have a higher average number of calls), it provides more evidence for that variable's importance.

4. Identifying Potential Predictors: The goal is to look for features or combinations of features that clearly separate churned customers from non-churned customers in the plots. These features are likely strong predictors of churn.

Preprocessing

encoding###

```
In [10]: #Convert categorical variables into numerical formats using one-hot encoding.
df[['international plan', 'voice mail plan']] = df[['international plan', 'voice ma
df = pd.get_dummies(df, columns=['state', 'area code'], dtype=int)
df.head()
```

```
Out[10]:
```

	account length	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls
0	128	382-4657	0	1	25	265.1	110	45.07	197.4	9
1	107	371-7191	0	1	26	161.6	123	27.47	195.5	10
2	137	358-1921	0	0	0	243.4	114	41.38	121.2	11
3	84	375-9999	1	0	0	299.4	71	50.90	61.9	8
4	75	330-6626	1	0	0	166.7	113	28.34	148.3	12

5 rows × 73 columns




```
In [11]: #Remove phone number, since it does not help in predicting churn
df = df.drop(columns=['phone number'])
```

```
In [12]: # Remove linear dependencies between features
numerical_features = [
    'account length', 'number vmail messages', 'total day minutes', 'total day call
    'total day charge', 'total eve minutes', 'total eve calls', 'total eve charge',
    'total night minutes', 'total night calls', 'total night charge', 'total intl m
    'total intl calls', 'total intl charge', 'customer service calls'
]

correlation_matrix = df[numerical_features].corr(method='spearman')
```

```
plt.figure(figsize=(10, 7))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix of Numeric Features')
plt.show()
```




Observation

- Low Correlation Across Most Features meaning no strong linear relationships exist
- Slight Negative Correlation with Customer Service Calls suggesting that customers making frequent service calls might not necessarily have high usage, indicating dissatisfaction rather than engagement.
- No Strong Predictive Feature Identified suggesting that feature engineering or non-linear modeling approaches may be necessary to extract meaningful patterns for churn prediction.

Churn Distribution

```
In [13]: plt.figure(figsize=(8, 6))
sns.countplot(x='churn', data=df)
```

```
plt.title('Churn Distribution')  
plt.xlabel('Churn (0: No, 1: Yes)')  
plt.ylabel('Number of Customers')  
plt.show()
```



The churn distribution shows an imbalance in the dataset. There are significantly more customers who did not churn **labeled as 0** than customers who did churn **labeled as 1**. This imbalance needs to be addressed during model training to prevent the model from being biased towards predicting the majority class **no churn**. Techniques like oversampling the minority class, undersampling the majority class, or using class weights during model training can help mitigate this issue and improve the model's ability to accurately predict churn.

```
In [14]: df = df.drop(columns=['total day charge', 'total eve charge', 'total night charge'],  
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3333 entries, 0 to 3332
```

```
Data columns (total 68 columns):
```

#	Column	Non-Null Count	Dtype
0	account length	3333 non-null	int64
1	international plan	3333 non-null	int64
2	voice mail plan	3333 non-null	int64
3	number vmail messages	3333 non-null	int64
4	total day minutes	3333 non-null	float64
5	total day calls	3333 non-null	int64
6	total eve minutes	3333 non-null	float64
7	total eve calls	3333 non-null	int64
8	total night minutes	3333 non-null	float64
9	total night calls	3333 non-null	int64
10	total intl minutes	3333 non-null	float64
11	total intl calls	3333 non-null	int64
12	customer service calls	3333 non-null	int64
13	churn	3333 non-null	int64
14	state_AK	3333 non-null	int64
15	state_AL	3333 non-null	int64
16	state_AR	3333 non-null	int64
17	state_AZ	3333 non-null	int64
18	state_CA	3333 non-null	int64
19	state_CO	3333 non-null	int64
20	state_CT	3333 non-null	int64
21	state_DC	3333 non-null	int64
22	state_DE	3333 non-null	int64
23	state_FL	3333 non-null	int64
24	state_GA	3333 non-null	int64
25	state_HI	3333 non-null	int64
26	state_IA	3333 non-null	int64
27	state_ID	3333 non-null	int64
28	state_IL	3333 non-null	int64
29	state_IN	3333 non-null	int64
30	state_KS	3333 non-null	int64
31	state_KY	3333 non-null	int64
32	state_LA	3333 non-null	int64
33	state_MA	3333 non-null	int64
34	state_MD	3333 non-null	int64
35	state_ME	3333 non-null	int64
36	state_MI	3333 non-null	int64
37	state_MN	3333 non-null	int64
38	state_MO	3333 non-null	int64
39	state_MS	3333 non-null	int64
40	state_MT	3333 non-null	int64
41	state_NC	3333 non-null	int64
42	state_ND	3333 non-null	int64
43	state_NE	3333 non-null	int64
44	state_NH	3333 non-null	int64
45	state_NJ	3333 non-null	int64
46	state_NM	3333 non-null	int64
47	state_NV	3333 non-null	int64
48	state_NY	3333 non-null	int64
49	state_OH	3333 non-null	int64
50	state_OK	3333 non-null	int64

```

51 state_OR          3333 non-null    int64
52 state_PA          3333 non-null    int64
53 state_RI          3333 non-null    int64
54 state_SC          3333 non-null    int64
55 state_SD          3333 non-null    int64
56 state_TN          3333 non-null    int64
57 state_TX          3333 non-null    int64
58 state_UT          3333 non-null    int64
59 state_VA          3333 non-null    int64
60 state_VT          3333 non-null    int64
61 state_WA          3333 non-null    int64
62 state_WI          3333 non-null    int64
63 state_WV          3333 non-null    int64
64 state_WY          3333 non-null    int64
65 area code_408     3333 non-null    int64
66 area code_415     3333 non-null    int64
67 area code_510     3333 non-null    int64

```

dtypes: float64(4), int64(64)

memory usage: 1.7 MB

Scaling###

```

In [15]: #split your dataset into training and testing sets
from sklearn.model_selection import train_test_split
X = df.drop('churn', axis=1)
y = df['churn']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st

```

```

In [16]: #Standardize the training and testing sets
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
scaler.fit(X_train)

X_train_st = scaler.transform(X_train)
X_valid_st = scaler.transform(X_test)

```

```

In [17]: # Convert scaled arrays back to DataFrames for easier inspection
X_train_scaled_df = pd.DataFrame(X_train_st, columns=X.columns)
X_test_scaled_df = pd.DataFrame(X_valid_st, columns=X.columns)

print("Training Data Mean:\n", X_train_scaled_df.mean())
print("\nTraining Data Standard Deviation:\n", X_train_scaled_df.std())

print("\nTesting Data Mean:\n", X_test_scaled_df.mean())
print("\nTesting Data Standard Deviation:\n", X_test_scaled_df.std())

```

Training Data Mean:

account length	-1.421654e-17
international plan	-4.975789e-17
voice mail plan	-6.184195e-17
number vmail messages	-4.620376e-17
total day minutes	1.610023e-16

...

state_WV	-2.061398e-17
state_WY	-1.421654e-17
area code_408	5.686617e-18
area code_415	-6.895023e-17
area code_510	-2.843308e-17

Length: 67, dtype: float64

Training Data Standard Deviation:

account length	1.0002
international plan	1.0002
voice mail plan	1.0002
number vmail messages	1.0002
total day minutes	1.0002

...

state_WV	1.0002
state_WY	1.0002
area code_408	1.0002
area code_415	1.0002
area code_510	1.0002

Length: 67, dtype: float64

Testing Data Mean:

account length	0.011575
international plan	-0.083161
voice mail plan	0.055053
number vmail messages	0.050326
total day minutes	-0.010897

...

state_WV	0.013575
state_WY	-0.023687
area code_408	0.056909
area code_415	-0.057964
area code_510	0.010369

Length: 67, dtype: float64

Testing Data Standard Deviation:

account length	0.989846
international plan	0.881556
voice mail plan	1.027189
number vmail messages	1.020775
total day minutes	0.981738

...

state_WV	1.036484
state_WY	0.923665
area code_408	1.032190
area code_415	0.999359
area code_510	1.006490

Length: 67, dtype: float64

Modeling

Regression

logistic regression

```
In [20]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

# Make predictions on the test set
y_pred = logreg.predict(X_valid_st)

# Evaluate the model
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))


accuracy = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy: {accuracy}")
print(f"Recall: {recall}")
print(f"Precision: {precision}")
print(f"F1-score: {f1}")
```

	precision	recall	f1-score	support
0	0.88	0.96	0.92	719
1	0.47	0.22	0.30	115
accuracy			0.86	834
macro avg	0.68	0.59	0.61	834
weighted avg	0.83	0.86	0.84	834

```
[[691 28]
 [ 90 25]]
Accuracy: 0.8585131894484412
Recall: 0.21739130434782608
Precision: 0.4716981132075472
F1-score: 0.2976190476190476
```

```
In [21]: # Plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```



Classification

Decision tree

```
In [22]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
from sklearn.metrics import ConfusionMatrixDisplay

clf = DecisionTreeClassifier(random_state=3)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)


print(f"Accuracy: {accuracy}")
print(f"Recall: {recall}")
print(f"Precision: {precision}")
print(f"F1-score: {f1}")
```

```
# Classification report
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9028776978417267
 Recall: 0.6869565217391305
 Precision: 0.6370967741935484
 F1-score: 0.6610878661087866

	precision	recall	f1-score	support
0	0.95	0.94	0.94	719
1	0.64	0.69	0.66	115
accuracy			0.90	834
macro avg	0.79	0.81	0.80	834
weighted avg	0.91	0.90	0.90	834

```
In [23]: cm = confusion_matrix(y_test, y_pred, labels=clf.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=clf.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for Decision Tree Classifier')
plt.show()
```



random forest####

```
In [24]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report
rf = RandomForestClassifier(random_state=3)
```

```

rf.fit(X_train_st, y_train)
y_pred = rf.predict(X_valid_st)
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy: {accuracy}")
print(f"Recall: {recall}")
print(f"Precision: {precision}")
print(f"F1-score: {f1}")

# Classification report
print(classification_report(y_test, y_pred))

```

Accuracy: 0.9328537170263789

Recall: 0.5478260869565217

Precision: 0.9402985074626866


F1-score: 0.6923076923076923

	precision	recall	f1-score	support
0	0.93	0.99	0.96	719
1	0.94	0.55	0.69	115
accuracy			0.93	834
macro avg	0.94	0.77	0.83	834
weighted avg	0.93	0.93	0.93	834

```

In [25]: cm = confusion_matrix(y_test, y_pred, labels=rf.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=rf.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for Random Forest Classifier')
plt.show()

```



K-NN model

```
In [26]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_st, y_train)
y_pred = knn.predict(X_valid_st)
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)


print(f"Accuracy: {accuracy}")
print(f"Recall: {recall}")
print(f"Precision: {precision}")
print(f"F1-score: {f1}")

# Classification report
print(classification_report(y_test, y_pred))
```


Accuracy: 0.8537170263788969
 Recall: 0.06086956521739131
 Precision: 0.3333333333333333
 F1-score: 0.10294117647058823

	precision	recall	f1-score	support
0	0.87	0.98	0.92	719
1	0.33	0.06	0.10	115
accuracy			0.85	834
macro avg	0.60	0.52	0.51	834
weighted avg	0.79	0.85	0.81	834

```
In [27]: cm = confusion_matrix(y_test, y_pred, labels=knn.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=knn.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for KNN Classifier')
plt.show()
```



SVM

```
In [28]: from sklearn.svm import SVC
from sklearn.metrics import confusion_matrix, classification_report, ConfusionMatrixDisplay
svm = SVC(random_state=3)
svm.fit(X_train_st, y_train)
y_pred = svm.predict(X_valid_st)
# Evaluate the model
```

```

accuracy = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy: {accuracy}")
print(f"Recall: {recall}")
print(f"Precision: {precision}")
print(f"F1-score: {f1}")

# Classification report
print(classification_report(y_test, y_pred))

```

Accuracy: 0.8669064748201439

Recall: 0.043478260869565216

Precision: 0.8333333333333334


F1-score: 0.08264462809917356

	precision	recall	f1-score	support
0	0.87	1.00	0.93	719
1	0.83	0.04	0.08	115
accuracy			0.87	834
macro avg	0.85	0.52	0.51	834
weighted avg	0.86	0.87	0.81	834

```

In [29]: cm = confusion_matrix(y_test, y_pred, labels=svm.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=svm.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for SVM Classifier')
plt.show()

```



Neural network

```
In [30]: from sklearn.neural_network import MLPClassifier

NeuralNet = MLPClassifier(random_state=5, max_iter=1000)
NeuralNet.fit(X_train_st, y_train)
y_pred = NeuralNet.predict(X_valid_st)

df_scores = pd.DataFrame({'Neural network': [f1_score(y_pred, y_test), accuracy_score(y_pred, y_test)]})
df_scores.T
```

```
Out[30]:
```


	f1_score	accuracy_score	recall_score
Neural network	0.494845	0.882494	0.607595

```
In [34]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

def show_confusion_matrix(y_true, y_pred):

    cm = confusion_matrix(y_true, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm)
    disp.plot(cmap=plt.cm.Blues)
    plt.show()
```

```
show_confusion_matrix(y_pred, y_test)
```



Naive Bayes

```
In [35]: from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report
nb = GaussianNB()
nb.fit(X_train_st, y_train)
y_pred = nb.predict(X_valid_st)
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)


print(f"Accuracy: {accuracy}")
print(f"Recall: {recall}")
print(f"Precision: {precision}")
print(f"F1-score: {f1}")

# Classification report
print(classification_report(y_test, y_pred))
```

Accuracy: 0.5779376498800959
 Recall: 0.591304347826087
 Precision: 0.18230563002680966
 F1-score: 0.2786885245901639

	precision	recall	f1-score	support
0	0.90	0.58	0.70	719
1	0.18	0.59	0.28	115
accuracy			0.58	834
macro avg	0.54	0.58	0.49	834
weighted avg	0.80	0.58	0.64	834

```
In [36]: cm = confusion_matrix(y_test, y_pred, labels=nb.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=nb.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for Naive Bayes Classifier')
plt.show()
```



Hyperparameter Tuning

tuned decision tree model

```
In [37]: from sklearn.model_selection import GridSearchCV
```

```

# Define the parameter grid for hyperparameter tuning
param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Create a GridSearchCV object
grid_search = GridSearchCV(estimator=clf, param_grid=param_grid, cv=5, scoring='f1')

# Fit the GridSearchCV object to the training data
grid_search.fit(X_train, y_train)

# Get the best hyperparameters and the best model
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_

print(f"Best hyperparameters: {best_params}")

# Evaluate the best model on the test data
y_pred = best_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy (tuned): {accuracy}")
print(f"Recall (tuned): {recall}")
print(f"Precision (tuned): {precision}")
print(f"F1-score (tuned): {f1}")

print(classification_report(y_test, y_pred))

```

```

Best hyperparameters: {'criterion': 'entropy', 'max_depth': 10, 'min_samples_leaf':
2, 'min_samples_split': 10}
Accuracy (tuned): 0.9436450839328537
Recall (tuned): 0.7217391304347827
Precision (tuned): 0.8469387755102041
F1-score (tuned): 0.7793427230046949

```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	719
1	0.85	0.72	0.78	115
accuracy			0.94	834
macro avg	0.90	0.85	0.87	834
weighted avg	0.94	0.94	0.94	834

tuned random forest model

```

In [38]: from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score

```

```

# Define the parameter grid for hyperparameter tuning
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['sqrt', 'log2']
}

# Create a GridSearchCV object
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, scoring='f1')

# Fit the GridSearchCV object to the training data
grid_search.fit(X_train_st, y_train)

# Get the best hyperparameters and the best model
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_

print(f"Best hyperparameters: {best_params}")

# Evaluate the best model on the test data
y_pred = best_model.predict(X_valid_st)
accuracy = accuracy_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy (tuned): {accuracy}")
print(f"Recall (tuned): {recall}")
print(f"Precision (tuned): {precision}")
print(f"F1-score (tuned): {f1}")

print(classification_report(y_test, y_pred))

```

Best hyperparameters: {'max_depth': None, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 50}

Accuracy (tuned): 0.9304556354916067

Recall (tuned): 0.5043478260869565

Precision (tuned): 0.9830508474576272

F1-score (tuned): 0.6666666666666666

	precision	recall	f1-score	support
0	0.93	1.00	0.96	719
1	0.98	0.50	0.67	115
accuracy			0.93	834
macro avg	0.95	0.75	0.81	834
weighted avg	0.93	0.93	0.92	834

Model evaluation

In [39]: `from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score`
`from sklearn.model_selection import GridSearchCV`

```

from sklearn.naive_bayes import GaussianNB
import matplotlib.pyplot as plt

def evaluate_model(model, model_name, X_test, y_test):
    # Changed X_valid_ to X_test to match function parameter and existing variables
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    print(f"{model_name} Metrics:")
    print(f"Accuracy: {accuracy}")
    print(f"Recall: {recall}")
    print(f"Precision: {precision}")
    print(f"F1-score: {f1}")
    print(classification_report(y_test, y_pred))


    cm = confusion_matrix(y_test, y_pred, labels=model.classes_)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.classes_)
    disp.plot(cmap=plt.cm.Blues)
    plt.title(f'Confusion Matrix for {model_name}')
    plt.show()

# Evaluate all models
evaluate_model(logreg, "Logistic Regression", X_valid_st, y_test)
evaluate_model(best_model, "Decision Tree", X_test, y_test)
evaluate_model(rf, "Random Forest", X_valid_st, y_test)
evaluate_model(knn, "K-NN", X_valid_st, y_test)
evaluate_model(svm, "SVM", X_valid_st, y_test)
evaluate_model(NeuralNet, "Neural Network", X_valid_st, y_test)
evaluate_model(nb, "Naive Bayes", X_valid_st, y_test)

```

Logistic Regression Metrics:
Accuracy: 0.8585131894484412
Recall: 0.21739130434782608
Precision: 0.4716981132075472
F1-score: 0.2976190476190476

	precision	recall	f1-score	support
0	0.88	0.96	0.92	719
1	0.47	0.22	0.30	115
accuracy			0.86	834
macro avg	0.68	0.59	0.61	834
weighted avg	0.83	0.86	0.84	834



Decision Tree Metrics:


Accuracy: 0.3381294964028777

Recall: 0.8695652173913043

Precision: 0.15698587127158556

F1-score: 0.26595744680851063

	precision	recall	f1-score	support
0	0.92	0.25	0.40	719
1	0.16	0.87	0.27	115
accuracy			0.34	834
macro avg	0.54	0.56	0.33	834
weighted avg	0.82	0.34	0.38	834



Random Forest Metrics:


Accuracy: 0.9328537170263789

Recall: 0.5478260869565217

Precision: 0.9402985074626866

F1-score: 0.6923076923076923

	precision	recall	f1-score	support
0	0.93	0.99	0.96	719
1	0.94	0.55	0.69	115
accuracy			0.93	834
macro avg	0.94	0.77	0.83	834
weighted avg	0.93	0.93	0.93	834



K-NN Metrics:


Accuracy: 0.8537170263788969

Recall: 0.06086956521739131

Precision: 0.3333333333333333

F1-score: 0.10294117647058823

	precision	recall	f1-score	support
0	0.87	0.98	0.92	719
1	0.33	0.06	0.10	115
accuracy			0.85	834
macro avg	0.60	0.52	0.51	834
weighted avg	0.79	0.85	0.81	834



SVM Metrics:


Accuracy: 0.8669064748201439

Recall: 0.043478260869565216

Precision: 0.8333333333333334

F1-score: 0.08264462809917356

	precision	recall	f1-score	support
0	0.87	1.00	0.93	719
1	0.83	0.04	0.08	115
accuracy			0.87	834
macro avg	0.85	0.52	0.51	834
weighted avg	0.86	0.87	0.81	834



Neural Network Metrics:


Accuracy: 0.882494004796163

Recall: 0.41739130434782606

Precision: 0.6075949367088608

F1-score: 0.4948453608247423

	precision	recall	f1-score	support
0	0.91	0.96	0.93	719
1	0.61	0.42	0.49	115
accuracy			0.88	834
macro avg	0.76	0.69	0.71	834
weighted avg	0.87	0.88	0.87	834



Naive Bayes Metrics:


Accuracy: 0.5779376498800959

Recall: 0.591304347826087

Precision: 0.18230563002680966

F1-score: 0.2786885245901639

	precision	recall	f1-score	support
0	0.90	0.58	0.70	719
1	0.18	0.59	0.28	115
accuracy			0.58	834
macro avg	0.54	0.58	0.49	834
weighted avg	0.80	0.58	0.64	834



Conclusion & Recommendation

Conclusion

Based on the evaluation of various classification models, including Logistic Regression, Decision Tree, Random Forest, K-NN, SVM, Neural Network, and Naive Bayes, the Random Forest classifier, after hyperparameter tuning, demonstrates superior performance in predicting customer churn for SyriaTel. This is evidenced by its high F1-score, accuracy, precision, and recall. While other models achieved reasonable performance, the Random Forest model's ability to balance precision and recall makes it the most suitable choice for this business problem. The tuned Decision Tree also performs well, but Random Forest provides a more robust prediction.

Recommendation

1. **Deploy the Tuned Random Forest Model:** Implement the tuned Random Forest model into SyriaTel's operational systems to predict customer churn in real-time. This allows for proactive customer retention strategies.

2. **Focus on Key Predictors:** Analyze the feature importances derived from the Random Forest model to identify the most significant drivers of customer churn. This will help focus retention efforts on the most impactful factors. Further investigation into these factors might reveal opportunities for service improvement or targeted marketing campaigns.
3. **Regular Model Updates:** Customer behavior and market dynamics can shift over time. Regularly retrain and evaluate the churn prediction model using updated data to ensure its continued accuracy and effectiveness.
4. **Implement a Customer Retention Strategy:** Based on the model's predictions, develop and implement targeted retention strategies for at-risk customers. This could include personalized offers, loyalty programs, improved customer service, or proactive communication.
5. **Consider Ensemble Methods:** Explore other ensemble methods or stacking techniques that combine the strengths of multiple models to potentially further improve predictive accuracy.
6. **Monitor Model Performance:** Continuously monitor the model's performance in a production environment and re-evaluate its effectiveness regularly.
7. **Explore External Data Sources:** Consider incorporating external data sources (if available) such as economic indicators, competitor activity, or market trends to enhance the model's predictive power.

In []: